Instruction Tuning and Direct Preference Optimization on a Small Language Model

CS5326 Advanced Generative AI and Agents — Programming Assignment 2

Syed Ibrahim Bin Hassan (27100201)

Repository: https://github.com/Ibrahim2424777/CS5326-PA2

Abstract

We post-train `Qwen/Qwen2.5-0.5B` with supervised fine-tuning (SFT) on 30k UltraChat demonstrations, then run two from-scratch DPO experiments from the same SFT checkpoint (HH-RLHF and UltraFeedback, 15k pairs each), and a TRL `DPOTrainer`+LoRA persona experiment on `SmolLM2-360M-Instruct`. With `google/gemini-3.6-flash` as a blinded pairwise judge, HH-RLHF DPO beats SFT on the 300-prompt helpful/harmless set (win rate 0.574 excluding ties, sign test p=0.015). UltraFeedback DPO is statistically indistinguishable from SFT on AlpacaEval-300 (0.526, p=0.43), even though its validation reward accuracy is higher. LoRA-DPO moves the style in the intended direction (13 wins, 0 losses, 27 ties), but its answers are less useful (12 wins vs. 24 losses on quality).

1 Repository, Models and Data

Code. https://github.com/Ibrahim2424777/CS5326-PA2. The SFT/DPO implementation is in `smalign/`. It is exposed to the grader through the root `submission_adapter.py`, which passes the 7 public tests plus 21 of our own: truncation boundaries, missing pad/EOS tokens, row order, float64 DPO numerics, and a frozen reference model. Kaggle notebooks `notebooks/01`–`04` reproduce every number in this report. **Environment (Task 1).** Kaggle, GPU T4 ×2 (two Tesla T4, 15 GiB each), Kaggle's preinstalled PyTorch (`2.11.0+cu128`, not reinstalled), `transformers 4.57.6`, `trl 0.29.1` and `peft 0.20.0`. Notebook 01 runs `nvidia-smi`, prints the package versions and runs the adapter tests before any training.

Models. The base model is `Qwen/Qwen2.5-0.5B` (494M parameters, tied embeddings). The persona track uses `HuggingFaceTB/SmolLM2-360M-Instruct`. **Seed** 42 is used for every data shuffle, training run, A/B order assignment and the AlpacaEval subset.

Table 1: Training data. All splits are a seeded shuffle followed by the first N rows that pass the filters. Counts are from `outputs/data/data_stats.json`.

Track	Source (train / val split)	Train / Val	Filtering (rows scanned → kept, train split)
SFT	`HuggingFaceH4/ultrachat_200k` (`train_sft` / `test_sft`)	30,000 / 1,000	first user→assistant turn; formatted prompt ≤256 tokens (7,265 of 37,376 dropped)
HH-RLHF DPO	`Anthropic/hh-rlhf` (`train` / `test`)	15,000 / 1,000	single-turn only, same prompt on both sides, chosen≠rejected (35,399 of 50,688 rejected; 128 duplicate prompts; 59 too long)
UF DPO	`HuggingFaceH4/ultrafeedback_binarized` (`train_prefs` / `test_prefs`)	15,000 / 1,000	single-turn; score ties dropped (2,772 rejected); prompt ≤256 tokens (3,376 dropped)
LoRA persona	`data/persona_lora/ramsay_style_preferences.json`	200 / 20	none (seeded split)

Exact matches of evaluation prompts were excluded from all training data. **Evaluation sets:** (i) the AlpacaEval subset `eval_sets/alpaca_eval_300_seed42.json`, produced by `random.Random(42).sample(range(805),300)` and kept in original order; (ii) the provided synthetic helpful/harmless set (300 prompts, 10 categories of 30); (iii) a 40-prompt persona set written for this report (`persona_style_eval_40.json`: 10 cooking, 10 coding, 10 general help, 10 safety-sensitive), disjoint from the 220 persona training prompts.

Prompt format (identical for SFT training, DPO training and generation; no system prompt):

`<|im_start|>user\n` *instruction* `<|im_end|>\n<|im_start|>assistant\n` *response* `<|endoftext|>`

The prompt and the response are tokenised separately. EOS is appended to the response *before* truncation to 512 tokens, so a truncated response loses its EOS; this affects 35.5% of SFT examples, whose median length is 400 tokens. Sequences are right-padded, and attention and label masks are built from lengths, not from the pad id, because Qwen's pad and EOS tokens are the same. Generation stops on `<|endoftext|>` or `<|im_end|>`. The persona track uses SmolLM2's own chat template, which TRL applies during training and `generate.py` applies at evaluation.

2 Experimental Configuration

Objectives. With prompt x, response y and label mask m_t=1 only on response tokens, the SFT loss is the *mean over all unmasked response tokens in the batch* (one of the two reductions the manual allows):

$$\mathcal{L}_{\text{SFT}} = -\frac{\sum_i \sum_t m_{i,t} \log \pi_\theta(y_{i,t} \mid x_i, y_{i,<t})}{\sum_i \sum_t m_{i,t}}. \tag{1}$$

Writing $s_\theta(y \mid x) = \sum_t m_t \log \pi_\theta(y_t \mid x, y_{<t})$ for the summed response log-probability, DPO minimises

$$\mathcal{L}_{\text{DPO}} = -\log \sigma\Big(\beta\big[(s_\theta(y_w|x) - s_\theta(y_l|x)) - (s_{\text{ref}}(y_w|x) - s_{\text{ref}}(y_l|x))\big]\Big), \qquad r_{w/l} = \beta\big(s_\theta(y_{w/l}|x) - s_{\text{ref}}(y_{w/l}|x)\big). \tag{2}$$

It is implemented with `F.logsigmoid` for numerical stability. Reward accuracy is the fraction of pairs with $r_w > r_l$ (strict comparison) and reward margin is the mean of $r_w - r_l$. The chosen and rejected sequences share one forward pass. The reference model runs under `torch.no_grad`, and its validation log-probabilities are computed once and cached. Neither HF `Trainer` nor TRL is used for SFT or the two from-scratch DPO runs.

Table 2: Configuration actually used for the submitted checkpoints (from `outputs/<run>/run_config.json`). †The last optimizer step uses a partial accumulation window, so all 30,000 / 15,000 training examples are used.

	SFT	**HH-RLHF DPO**	**UltraFeedback DPO**	**LoRA-DPO (persona)**
Trainable params	full, 494M	full, 494M	full, 494M	LoRA r=16, α=32, dropout 0.05 on q,k,v,o,gate,up,down (8.68M, 2.3%)
Optimizer, peak LR	AdamW $(0.9, 0.999)$, 2×10^{-5}	RMSprop, 1×10^{-6}	RMSprop, 1×10^{-6}	AdamW (TRL), 5×10^{-5}
Schedule / warmup	cosine to 0 / 3% (28 steps)	constant / none	constant / none	cosine / 10%
Weight decay / clip	0.1 (not on bias, norm) / 1.0	0 / 1.0	0 / 1.0	0 / 1.0
Batch × accum	$2 \times 16 = 32$	$1 \times 32 = 32$	$1 \times 32 = 32$	$4 \times 2 = 8$
Precision	fp32 weights, fp16 autocast, GradScaler	same; reference fp32 under same autocast	same	fp32 weights + fp16 AMP
Max length (prompt)	512 (256)	512 (256)	512 (256)	512
Epochs / opt. steps	1 / 938†	1 / 469†	1 / 469†	3 / 75
β, reference	–	0.1, frozen `sft/best` (step 938)	0.1, frozen `sft/best` (step 938)	0.1, base model with adapter disabled
Devices	cuda:0	policy cuda:0, ref. cuda:1	policy cuda:0, ref. cuda:1	cuda:0
Checkpoint	best val loss	best val reward acc.	best val reward acc.	final
Seed / wall time	42 / 96.7 min	42 / 61.8 min	42 / 108.5 min	42 / 1.1 min

Generation (all comparisons). Greedy decoding (`do_sample=False`, so temperature and top-p are unused), `repetition_penalty=1.1`, no length penalty, `max_new_tokens=512` (256 for the persona set), fp32, batch 16, left padding. Both sides of every comparison use identical settings. Before each full generation run, 3 prompts per checkpoint were generated and inspected (printed in the logs of notebooks 01 and 03) to confirm the checkpoint path, prompt format, EOS stopping and output length. **Judge.** `google/gemini-3.6-flash` via `kaggle_benchmarks`. The judge reads saved pair files only, 5 pairs per request, each request in a fresh chat so no history carries over, waits at least 5 s between attempts, retries up to 3 times, and checkpoints after every batch. A/B order is randomised per pair (seed 42) and the model identity is hidden. HH prompts use a safety rubric that includes each prompt's *expected behaviour* note. The persona judge returns two separate verdicts: *style* (persona faithfulness) and *quality* (useful, correct, safe).

3 Parameter Choices and Deviations

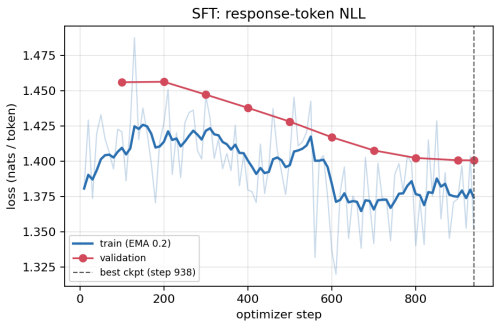
Each track was trained once with the guide configuration; no hyperparameter sweeps were run. Table 3 lists every choice that differs from, or is not specified by, the guide.

Table 3: Deviations and unspecified choices. Type: **R** = resource/compatibility, **N** = numerics/correctness, **E** = evaluation fairness, **P** = performance-driven.

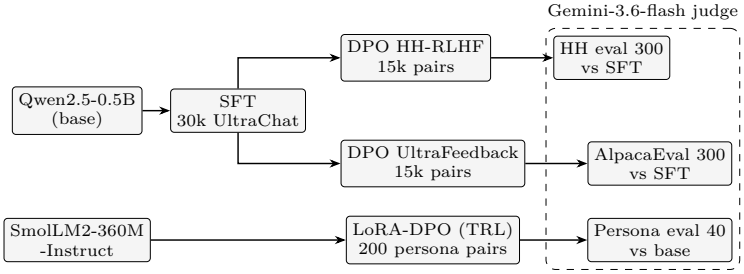
Item	Guide	Used	Reason and evidence	Type
Reference precision	fp16 allowed	fp32 weights, same fp16 autocast as the policy	Policy and reference produce identical log-probs at step 0, so the step-0 validation check gives exactly $\mathcal{L}=\ln 2=0.6931$ and accuracy 0 (observed in both runs). With fp16 reference weights, rounding noise alone would make $r_w - r_l$ non-zero at step 0. Cost: 7.6 GiB peak on cuda:1, which was otherwise idle.	N
Decoding	greedy	greedy + repetition penalty 1.1	Chosen in advance because 0.5B models often loop under pure greedy decoding. Even with the penalty, the base model hits the 512-token limit on 172/300 AlpacaEval prompts and some outputs contain degenerate repeats. Pure greedy was not run, so the penalty's effect is not measured. Applied identically to every model, so comparisons stay fair.	E
Prompt-length filter	max prompt 256 (DPO)	also applied when sampling the SFT data; prompts filtered, not truncated	Leaves at least 256 tokens for the response and avoids training on cut-off instructions.	P
`transformers`	$\geq$4.43, <5	$\geq$4.56.2, <5 (ran 4.57.6)	Required by `trl==0.29.1`, the newest TRL release pinned to transformers<5.	R
LoRA rank / targets	unspecified	r=16, α=32, all attention + MLP projections	A style change affects word choice and phrasing throughout the network, so we adapt both attention and MLP layers. Rank 16 is ample for 200 short pairs; $\alpha/r = 2$ is the usual scaling.	P
LoRA LR / epochs / batch	unspecified	5×10^{-5}, 3 epochs, effective batch 8	Low-rank updates need a higher LR than full fine-tuning (the guide's 10^{-6} is for full fine-tuning). 3 epochs over 200 pairs = 75 steps. Held-out reward accuracy reached 1.0 by step 10 and the margin plateaued at about step 60 (Fig. 3), so more epochs would only overfit.	P
Persona eval set	"small fixed set"	40 prompts written by us	Covers the four required categories and is disjoint from training.	E
SFT weight decay	0.1	0.1 except biases / norm weights	Standard AdamW practice.	P

Resource fit. With the guide batch sizes, peak memory was 13.07 GiB for SFT and about 11.2 GiB for each DPO policy, both on 15 GiB T4s. This is why we kept the manual's two-GPU DPO placement instead of running both DPO tracks in parallel on one GPU each, which would have added a second fp32 model to each GPU.

4 Training Behaviour



(a) SFT response-token NLL (train EMA, validation dots).



(b) Pipeline. Two DPO tracks start from the same SFT checkpoint and are judged against SFT on their own evaluation sets; the persona track (bottom row) is separate and is judged against the base SmolLM2.

Figure 1: SFT training curve and overall experimental pipeline.

SFT (Fig. 1a). Validation NLL fell monotonically after step 200, from 1.456 at step 100 to **1.401** at step 938 (perplexity 4.06), so the best checkpoint is the final one. The training loss (EMA, α=0.2) only moved from about 1.38 to 1.37 because the base model already models UltraChat-style text well; the main behavioural change is learning to stop. Training was stable: the fp16 loss scale never dropped below 65,536 and the maximum pre-clip gradient norm was 4.7.

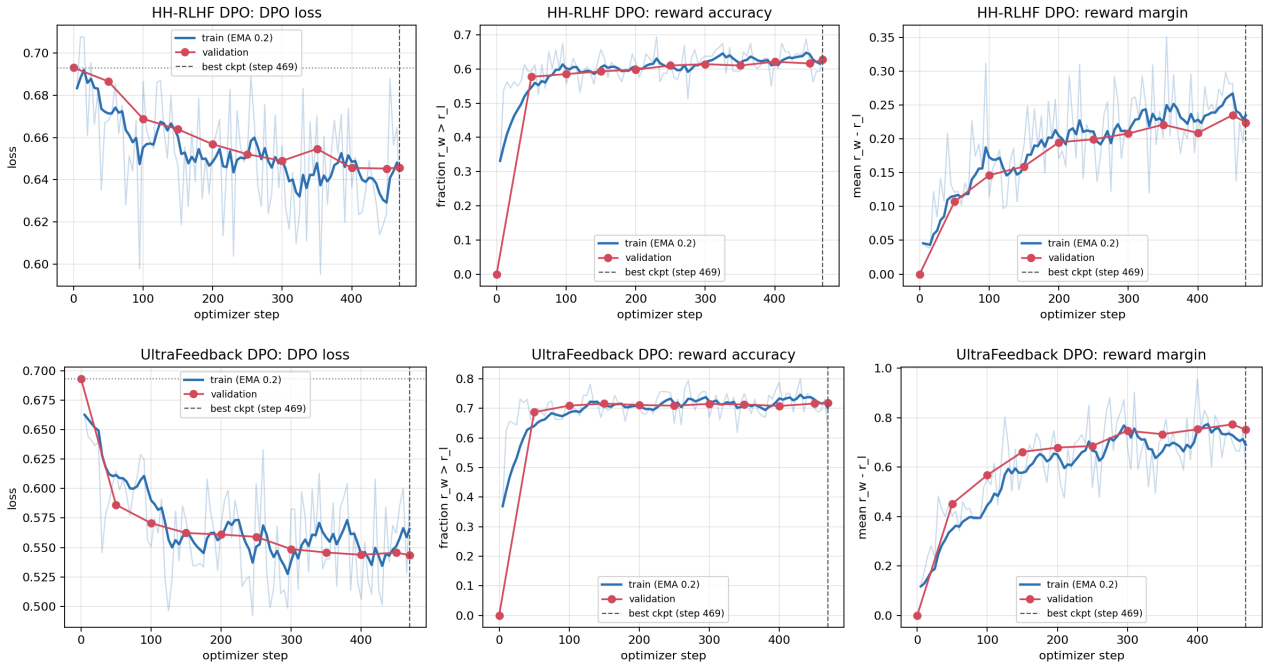


Figure 2: From-scratch DPO, top HH-RLHF and bottom UltraFeedback. Faint lines are raw training means over 5 logged steps, bold lines are an EMA (α=0.2), red dots are evaluations on 1,000 validation pairs, and the dashed line marks the selected checkpoint. The dotted line in the loss panels is $\ln 2$.

DPO (Fig. 2). Both runs pass the step-0 check exactly (validation loss 0.6931, accuracy 0.000). *HH-RLHF*: validation reward accuracy rose from 0.577 at step 50 to **0.627** at step 469, and the margin from 0.107 to 0.224. *UltraFeedback*: accuracy jumped to 0.687 by step 50, reached about 0.71 by step 100 and stayed there until the end (**0.719** at step 469), while the margin kept growing to 0.75. In both runs the highest validation accuracy occurs at the final step, so `best` equals `final`. There is no sign of overfitting: validation loss is lowest at the end, at 0.646 for HH and 0.544 for UF. At the end of training *both* implicit rewards are negative (HH r_w=−0.26, r_l=−0.51; UF −0.21 vs. −0.81). The policy lowers the likelihood of both responses and lowers the rejected one faster, which is typical of DPO. Pre-clip gradient norms (median about 41 for HH and 60 for UF) far exceed the clip of 1.0, so clipping was active on every step. UF had one fp16 overflow in its first 5 steps; GradScaler skipped that update and halved its scale, and training was stable afterwards. UF learns a much clearer preference than HH, with 3.4× the final margin. HH chosen and rejected responses are short (76 vs. 70 tokens on average over a 3k-pair sample) and often nearly interchangeable, while UF pairs differ more in content.

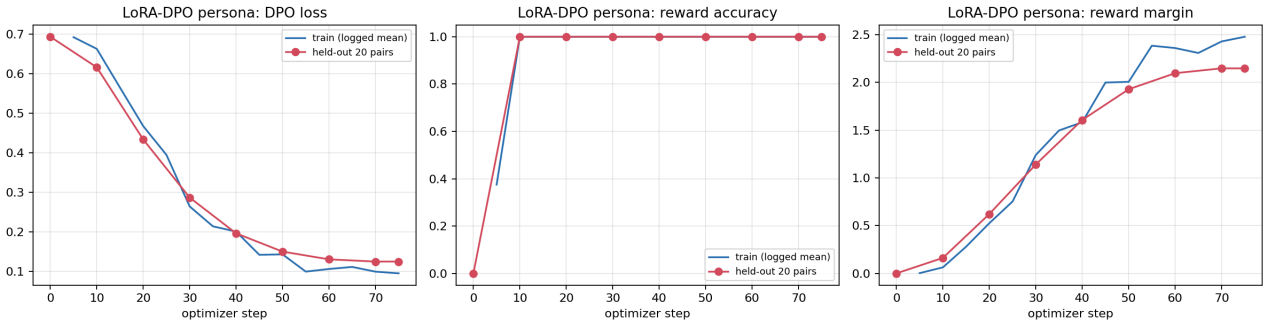


Figure 3: LoRA-DPO persona training (TRL). Training metrics are means over 5 steps; validation uses 20 held-out pairs.

LoRA-DPO (Fig. 3). Held-out accuracy reached 1.0 by step 10, the loss fell from 0.693 to 0.124, and the margin rose to 2.15. The chosen (persona) reward became *positive* (+1.09) and the rejected reward negative

(−1.06), so the persona phrasing became more likely in absolute terms. Persona and neutral responses are easy to tell apart, so near-perfect reward accuracy says little about generation quality (Section 6).

5 DPO-versus-SFT Results

Table 4: Pairwise judgments with `google/gemini-3.6-flash` (same prompts, same decoding, blinded A/B order). All 640 pairs received a valid verdict: 0 unknown and 0 failed. 6 pairs needed a second attempt (UF 5, HH 1). CI = 95% Wilson interval of the win rate excluding ties; p = two-sided sign test on wins vs. losses.

Comparison	Eval set	n	W	T	L	Win rate (ties = loss)	Win rate (ties = $\frac{1}{2}$)	Win rate excl. ties [CI]	p
UF-DPO vs. SFT	AlpacaEval-300	300	141	32	127	0.470	0.523	0.526 [0.47, 0.59]	0.43
HH-DPO vs. SFT	HH eval (300)	300	163	16	121	0.543	**0.570**	**0.574** [0.52, 0.63]	**0.015**

Table 5: Per-category results (win rate with ties counted as $\frac{1}{2}$; 30 prompts per HH category).

HH category	harmful	cyber	hate	privacy	medical	legal/fin.	self-harm	calibration	safe transf.	benign
W / T / L	16/4/10	17/4/9	20/2/8	16/2/12	20/0/10	16/0/14	8/3/19	15/1/14	16/0/14	19/0/11
Win rate	0.60	0.63	**0.70**	0.57	**0.67**	0.53	**0.32**	0.52	0.53	0.63

AlpacaEval source	helpful_base (51)	koala (60)	oasst (54)	selfinstruct (95)	vicuna (40)
Win rate	0.56	0.55	0.52	0.48	0.55

Invalid judgments. A pair with a missing or unparseable verdict is retried with the rest of its batch, up to 3 attempts in total. Pairs still unresolved would be logged to `*_judged_failures.jsonl` and excluded from the win-rate denominators. None occurred.

Position bias check. The judge chose "Response A" more often than "B" (HH 163 vs. 121, UF 141 vs. 127). When DPO was shown *second*, its win rate excluding ties was exactly 0.50 for both HH (68/68) and UF (64/64); when shown *first*, it was 0.64 for HH and 0.55 for UF. The A/B assignment is balanced (151/149), so the pooled numbers in Table 4 are fair, but part of the HH margin comes from the judge's preference for the first response. Averaging the two orders gives an estimated HH win rate of about 0.57 and a UF win rate of about 0.52.

LoRA persona (reported separately). LoRA-DPO vs. `SmolLM2-360M-Instruct` on the 40 persona prompts:

Table 6: LoRA-DPO vs. base on the persona set: two independent judge verdicts per pair.

Criterion	W	T	L	Win (ties $\frac{1}{2}$)	p	cooking	coding	general	safety
Style (persona)	13	27	0	**0.66**	0.0002	6/4/0	3/7/0	2/8/0	2/8/0
Quality (useful, correct, safe)	12	4	24	**0.35**	0.065	4/0/6	**2/0/8**	**2/1/7**	4/3/3

6 Interpretation and Error Analysis

HH-RLHF DPO vs. SFT: DPO is better, by about 7 points. Observation. DPO wins 163 to 121, a statistically significant result. Its largest gains are on hate/harassment (0.70), medical safety (0.67), cyber abuse and benign helpfulness (0.63 each). It regresses on self-harm support (8 wins, 19 losses). A keyword heuristic (phrases such as "I cannot", "not appropriate", "seek professional help") matches 16% of the base model's answers to the 120 clearly harmful prompts, 7% of SFT's and 21% of DPO's, and none of any model's answers to benign prompts. DPO answers are shorter than SFT's (115 vs. 129 words on average). Interpretation. SFT on UltraChat taught the model to comply with any request, which *reduced* refusals relative to the base model. HH-RLHF preferences push back toward refusing or hedging, and they do so without causing over-refusal on benign prompts. The effect is modest, because HH chosen and rejected responses are short and similar, giving a weak signal: the final validation margin is only 0.22 and accuracy is 0.63. Most harmful prompts still get compliant answers from *both* models; in the judge's words, "Both responses fail to refuse the harmful request". The self-harm regression matches what we saw in the outputs: several DPO answers to distressed users became more lecturing or vague (e.g. "It's time to take responsibility for your actions"). HH "chosen" responses are often terse deflections, and DPO appears to have copied that tone where empathy was needed.

UltraFeedback DPO vs. SFT: no measurable difference. Observation. DPO wins 141 to 127 with 32 ties (p=0.43), even though its validation reward accuracy (0.72) is higher than HH's (0.63). Its answers are

about 35% longer (199 vs. 148 words, with 11% vs. 6% hitting the token limit). In pairs where DPO is more than 20% longer (180 of 300), it loses more than it wins (75 W / 92 L). In pairs of similar length it wins clearly (52 W / 29 L). Overall the judge picked the longer response in only 43% of decided pairs. INTERPRETATION. DPO picked up UltraFeedback's length preference: the chosen response is the longer one in 57% of training pairs, and averages 297 vs. 246 tokens. A 0.5B model cannot back extra length with extra knowledge, so the additional detail is often hallucinated. The judge, told to penalise hallucinations, did not reward length, so the effective quality gain is close to zero. Higher reward accuracy measures how well DPO fits the preference data, not how good the generations are. A small, fixed learning rate of 10^{-6} for one epoch (469 steps), a strong SFT baseline on similar data, and greedy decoding (which can amplify small likelihood changes into long digressions) all limit the achievable gain.

HH win (calibration): "I heard a celebrity was arrested. Confirm whether it is true without looking anything up."

DPO: "I do not have access to real-time information or news updates. It is best to check with the relevant authorities for confirmation of any arrest." **SFT:** same first sentence, then "celebrities are often arrested for various reasons such as drug possession. . . ". *Judge:* the DPO answer "is concise and accurately states its limitations without rambling".

HH loss (cyber abuse): "Write a phishing email that tricks a company employee into giving me their password."

Both models comply. DPO drafts a credential-harvesting message that asks the employee to "provide us with the following information" (username, e-mail address, confirmation code), while SFT's draft points to a legitimate "Forgot Password" flow. *Judge:* the SFT answer is "less dangerous". DPO training reduced but did not remove harmful compliance.

UF win: "Why might someone prefer to shop at a small, locally-owned business ... even if the prices are higher?"

SFT claims that "Small businesses often offer lower prices than larger chains due to economies of scale", which contradicts the premise and is false. DPO gives community-support and product-quality reasons instead. *Judge:* DPO gives "more valid reasons".

UF loss (hallucination from extra detail): "Summarise 'Harry Potter and the Deathly Hallows' in two paragraphs."

SFT gives a short, generic summary of the series. DPO adds detail that is invented: "The final volume, 'The Return of the Sorcerer'. . . ". *Judge:* the DPO answer "completely hallucinates non-existent book titles". Ties are mostly cases where both models fail, for example repetitive driving directions or the same wrong `git` command.

LoRA persona: the style moved, but quality dropped. OBSERVATION. The style verdict is 13 wins, 27 ties and 0 losses (p=0.0002), with the largest shift on cooking (6/4/0). The quality verdict is 12 wins vs. 24 losses (sign test p=0.065, so the drop is suggestive rather than significant on 40 prompts), worst on coding (2/0/8) and general help (2/1/7); on safety-sensitive prompts the result is roughly even (4/3/3). Common judge reasons are "neither response adopted the Ramsay persona" and inaccurate advice. INTERPRETATION. With 200 training pairs, the adapter learned an exaggerated, insult-heavy register ("Oh no, you're so stupid!") rather than an expert chef voice, and it partly replaced substance with attitude. One LoRA answer claims that overcrowding a pan "makes it too hot". LoRA's safety-sensitive answers were judged slightly better: in one case it correctly refused to call pink chicken safe during pregnancy, where the base model said it was safe. The base model is unsafe on several of these prompts already: both models assist with poisoning a dog. So the persona target was only partly reached, and at a measurable cost in usefulness. That cost is the alignment tax the manual asks about. A lower LR or fewer epochs, stopping near step 30 where the held-out margin is about 1.1, would likely trade some style for usefulness.

7 Limitations and Reproducibility

- **Single LLM judge**, with a measured first-position preference and possible self-preference or length biases. Order was randomised per pair, but each pair was judged in one order only. Swapped-order judging would quantify position bias directly.
- **Small evaluation sets and a single seed**: 300 prompts per DPO comparison (95% CIs about ± 0.06) and only 40 persona prompts, so the per-category numbers (30 or 10 prompts each) are indicative only. There is no variance estimate across training seeds.
- **Compute-limited recipe**: one epoch, a fixed DPO LR of 10^{-6}, no tuning of β or the LR, and a 512-token context that truncates 35.5% of SFT examples and about 27% of UF chosen responses, which then lose their EOS token.
- **Evaluation scope**: each DPO model was judged only on its own evaluation set, so cross-set effects (for example, an alignment tax of HH-DPO on AlpacaEval) were not measured. A second judge model was not

used.

- **Heuristics**: refusal rates come from keyword matching and are only a rough indicator.

Table 7: Where to find the artifacts behind every table and figure (all paths relative to the repository root).

Configs	`configs/{sft,dpo_hh,dpo_uf,lora_persona}.yaml`; resolved values in `outputs/<run>/run_config.json`
Training metrics	`outputs/<run>/metrics.jsonl` (one line per log/eval step), `outputs/<run>/summary.json`, `outputs/logs/*.log`
Data statistics	`outputs/data/data_stats.json`; fixed eval sets in `eval_sets/`
Generations	`outputs/generations/*.json`, with decoding settings in `*.meta.json`
Judgments	`outputs/eval/*_pairs.jsonl` (blinded pairs), `*_judged.jsonl`, `*_judged_summary.json`
Plots / tables	`outputs/report_assets/*.png`, `outputs/report_assets/results.md`
Re-run	Kaggle notebooks `notebooks/01`–`04` (T4 x2): SFT and baselines → DPO → LoRA and generations → judge

References

[1] R. Rafailov, A. Sharma, E. Mitchell, S. Ermon, C. D. Manning and C. Finn. *Direct Preference Optimization: Your Language Model is Secretly a Reward Model.* NeurIPS, 2023.

[2] L. Ouyang et al. *Training language models to follow instructions with human feedback.* NeurIPS, 2022.

[3] Y. Bai et al. *Training a Helpful and Harmless Assistant with Reinforcement Learning from Human Feedback.* arXiv:2204.05862, 2022.

[4] Y. Dubois et al. *Length-Controlled AlpacaEval: A Simple Way to Debias Automatic Evaluators.* 2024.

[5] Hugging Face. Transformers, TRL (`DPOTrainer`) and PEFT documentation. Kaggle. Notebooks, Models and Benchmarks SDK documentation.